

CATHY BROWN ADVISORY

AI and Automation Maturity Framework for Finance Leaders

A practical guide to knowing where you are on the AI and automation journey, and exactly what to do next to move forward.

Cathy Brown, ACA

Founder, Cathy Brown Advisory
cathybrownadvisory.co.uk

Why this framework exists

Every finance leadership event asks the same question in different words: where do I even start with AI? There is no shortage of tools, only a shortage of a clear, honest sequence to follow.

This framework gives you that sequence, built specifically for finance leaders. Two things make it different from a typical tool list: it separates build capability (levels 1 to 5, what you can operate) from deployment posture (level 6, what risk you accept once a system can act on its own) rather than treating them as one ladder. And it treats data literacy and governance as foundational, not optional - a discipline that runs underneath every level rather than a stage you complete and move past.

Work through the levels in order, and use the free resources listed at each one to build genuine fluency, not just familiarity.

All links in this document were checked and confirmed valid as at 20 July 2026.

The framework at a glance

Levels 1 to 5: build capability, what you can operate and create.

Level 6: deployment posture, what risk you accept once a system can act on its own.

LEVEL 1

Chat-level LLMs

Claude, ChatGPT, Gemini, Perplexity, Grok - conversational use, one question and answer at a time



LEVEL 2

Projects and structured research

Organising AI around a body of work, e.g. NotebookLM, Claude Projects



LEVEL 3

AI embedded in your daily tools

Cowork, Copilot sidebar - AI sitting inside the work itself



LEVEL 4

No-code automation chains

Power Automate, Make, n8n - linking steps without writing code



LEVEL 5

Building bespoke tools

Splits into two paths depending on technical comfort



5A NO-CODE

No-code building

Lovable, Base44, Bolt, Copilot Studio: describe it in plain language, the platform builds it



5B CODE-ASSISTED

Code-assisted building

Claude Code, Cursor, GitHub Copilot agent mode: AI writes and debugs the code, you review it

LEVEL 6

Agentic systems

Copilot Cowork, Claude Cowork, Copilot Studio - split by oversight design

Data literacy and governance

Instead of being a stage to complete, this should run underneath every level on the ladder. It is a discipline to build from day one, and it matters more as you climb the ladder, not less – the more autonomy you hand to a system, the more it matters that you trusted it with the right information in the first place and checked its output before acting on it. In practice, that means staying alert to what is safe to share, aware of hallucination risk, and always verifying output before you act on it.

Two questions worth asking every time

- Before sharing anything: would I be comfortable with this information being read by someone outside the business?
- Before acting on anything: have I checked this is right, or am I assuming it is because it sounds confident?

Anthropic's prompt engineering course (see Level 2) touches on the output-verification side of this. The data-sharing discipline is more a habit you build through repetition than something a single course will install for you – but naming it explicitly, as this framework does, is the first step.

Level 1

Chat-level LLMs

This is where almost everyone starts, and where the data literacy and governance foundation needs to be established from day one, not added on later.

Fluency at this level means using a chat tool like Claude or ChatGPT for real, recurring work tasks rather than occasional curiosity. Draft an email. Summarise a long document before a meeting. Ask it to explain something you only half understand. The test is not whether you have tried it once - it is whether reaching for it has become a habit.

TRY THIS

Use Claude or ChatGPT for one real task every working day for two weeks. Drafting, summarising, or explaining - not just asking trivia questions.

Anthropic Academy's three foundational courses work well together and are all free:

[Claude 101](#) (Projects, Artifacts, basic prompt strategy);

[AI Fluency: Framework and Foundations](#) (how to work with AI as a thinking partner, not just a tool); and

[AI Capabilities and Limitations](#) (what AI can genuinely do versus where it confidently gets things wrong - the foundation of the data literacy and governance thread running through every level).

Level 2

Projects and structured research

The shift here is from a single question to a body of work. Instead of starting from zero each time, you give the AI persistent context so it understands the wider task.

This is where tools like NotebookLM and Claude Projects come in. The fluency test is whether you can set up a project around something real - a board pack you are drafting, a process you are documenting - and feel the difference that persistent context makes versus a single chat exchange.

TRY THIS

Set up a Claude Project or a NotebookLM notebook around one real, ongoing piece of work this month.

Work through Anthropic's free [interactive prompt engineering tutorial](#). Structuring what you tell the AI is what makes Projects and NotebookLM worth the extra setup.

Level 3

AI embedded in your daily tools

AI now sits inside the software you already use, rather than being a separate destination you have to remember to visit. Copilot and the Copilot sidebar in Word and Excel are the clearest examples.

One important distinction worth knowing if you use Microsoft 365: the standard Copilot sidebar in your everyday apps sits at this level. It is not the same as the fully autonomous Copilot Copilot product, which sits much higher up this ladder. Knowing which one you are using matters for how much you should trust its output unsupervised.

TRY THIS

If you are in the Microsoft ecosystem, complete Microsoft Learn's free [Copilot for Finance use case module](#). It walks through Excel COGS analysis, summarising meeting notes, evaluating an acquisition in Copilot Chat, and building an executive presentation.

Do the exercises against your own real numbers rather than the sample data, so you can judge the output against what you already know to be true.

Level 4

No-code automation chains

Tools like Power Automate and Make let you link steps together, so a process runs without you doing each part by hand, without writing code.

If you haven't already automated your data prep with Power Query or repetitive Excel tasks with Office Scripts, that's worth doing in parallel. Neither is an AI tool - they're deterministic, not intelligent - but they're often what feeds a Level 4 automation clean data in the first place.

This level has more depth inside it than the others, from a simple two-step automation through to calling external APIs directly and orchestrating a multi-step pipeline with conditional logic. Where Power Automate and Make cover the earlier stages well, n8n is worth knowing about for the later ones. It is self-hostable, gives full API access rather than relying on pre-built connectors, and includes a code step for JavaScript or Python - which is what lets it reach Stage 4 and beyond. It still sits at Level 4, not Level 5: even with that code step, it is orchestrating and connecting things that already exist, not generating a new piece of software.

STAGE 1 Single trigger, single action

Connecting two apps with one straightforward automation.

STAGE 2 Filters and routing logic

Adding conditions so the automation behaves differently depending on the data.

STAGE 3 Multiple data sources

Pulling from more than one place - a form, a spreadsheet, a shared drive.

STAGE 4 Calling external APIs directly

Constructing your own calls to a service like the Claude API, not just a pre-built connector.

STAGE 5 Multi-service orchestration

Several tools talking to each other in sequence, each handing off to the next.

STAGE 6 Error handling and production hardening

Retries and failure alerts so the pipeline can run unattended.

Level 4 (continued)

TRY THIS

Start small. Automate one repetitive task you do by hand, such as moving data between a form and a spreadsheet.

If you're using Power Automate, Microsoft Learn's free module on [using Copilot inside Power Automate](#) is a good next step. Microsoft also runs free [one-day Power Platform in a Day workshops](#), both live and self-paced, for a fuller grounding.

If you're using Make instead, [Make Academy](#) is free and self-paced - start with the Foundation course to build your first scenario.

If you're pushing into n8n for the later stages, the official [n8n Academy](#) is free, and its N8N101 Essentials course covers building your first workflows.

Level 5

Building bespoke tools

This level is genuinely different in kind from Level 4, not just more complex. Level 4 orchestrates tools and APIs that already exist. Level 5 generates a brand-new piece of software from a description of what you want.

It splits into two paths depending on technical comfort, and neither is more advanced than the other - they are different routes to the same outcome.

The no-code path uses tools like Lovable, Base44, and Bolt, where you describe what you want to build in plain language and the tool generates a working application from that description. You are directing the build rather than writing it. This path suits finance leaders who want to create client-facing tools, internal dashboards, or diagnostic products without involving a developer. If you're in the Microsoft ecosystem, Copilot Studio offers a similar no-code route - provided the agent you build stays purely conversational, answering when asked rather than acting on its own. Add an autonomous trigger and the same build moves to Level 6 instead (more on that distinction when we get there).

The code-assisted path uses tools like Claude Code, Cursor, and GitHub Copilot's agent mode (the natural choice if you're already in the Microsoft ecosystem), where AI writes, edits, and debugs code on your behalf within a development environment. This path genuinely requires a degree of prior coding experience – not because you will be writing code from scratch, but because you need to be able to read what the AI produces, understand broadly what it is doing, and spot when something has gone wrong. Without that foundation, you are approving output you cannot evaluate, which removes the main safeguard. This path suits people who have coded before and want to build something with deeper technical control, or who are willing to invest time building that foundation first.

In both cases, the output is a tool someone else will use. That's the defining characteristic of Level 5 - you are no longer just using AI, you are building with it.

Level 5 (continued)

5a No-code building

Tools like Base44, Lovable, and Bolt, where you describe what you want and the platform builds it without you touching code. Base44 leans toward internal business tools and diagnostics, Lovable toward full customer-facing web applications, and Bolt toward projects where you may want to hand off clean, exportable code later. If you're in the Microsoft ecosystem, Copilot Studio offers a similar no-code route - but it only stays at this level while the agent is purely conversational. Add an autonomous trigger and the same build becomes a Level 6 system instead (see the note on Copilot, Studio, and Cowork in Level 6).

5b Code-assisted building

Tools like Claude Code, Cursor, and GitHub Copilot's agent mode, where AI writes, edits, and debugs code on your behalf and you review it. Of the three, Claude Code is the gentlest entry point for non-coders: its browser and desktop versions need no installation or development environment, and you can describe a task in plain language and review the output without touching a terminal. Cursor and GitHub Copilot's agent mode both sit inside a developer IDE (VS Code or similar), so the interface itself assumes some familiarity with code before you start. Genuine technical comfort with any of these unlocks more flexibility, but the on-ramp into them is not the same.

A note on Claude Code: it genuinely has agentic mechanics; it plans and executes multi-step work with varying degrees of autonomy. It sits here rather than at Level 6 because of what it is for, not how it behaves. Level 5 tools build something new. Level 6 tools run an ongoing business process on their own. Claude Code's agentic behaviour is in service of writing and editing code, the output is a tool, not a deployed autonomous worker.

TRY THIS

For 5a: pick one real, low-stakes problem - a small calculator or internal tool - and try building it in Base44 or Lovable from a plain-English description.

For 5b: the lowest-friction starting point is Claude Code in your browser or desktop app - no installation needed. Describe a small, real task in plain language (a calculator, a simple data formatter, a utility you need) and let it plan and build a first draft. Your job is to review what it produces and steer it, not write a line of code yourself. If you find that comfortable and want to go further, Cursor and GitHub Copilot's agent mode offer more control - but both sit inside a developer IDE (Integrated Development Environment), so expect a steeper setup before you get started.

For those who already have some coding background and want to go deeper, Anthropic Academy's free [Claude Code in Action course](#) is a good place to start.

Level 6

Agentic systems

Agentic systems, whether built in Copilot Studio, Claude Cowork, custom workflows, or finance platforms such as BlackLine and Datarails, can plan and execute multi-step work with limited human intervention. This makes the second dimension of the framework, deployment posture, the key governance question, because it is separate from how the system was technically built.

Deployment posture is about how much autonomous action you are comfortable allowing a system to take without human approval at every step. A Level 4 automation runs a defined sequence. A Level 6 agent makes decisions about what to do next based on what it encounters. That is a materially different governance proposition.

For finance leaders, the questions at this level are specific and important. What is the agent authorised to do, and what requires human sign-off? An agent that can read data is very different from one that can write to a system, send communications, or trigger payments. How is the agent's output being verified, and what controls exist if it makes a poor decision? The hallucination and fabrication risks discussed earlier do not disappear at Level 6 - they increase, because the system is acting on its own outputs without a human reviewing each step. Who is accountable when something goes wrong? In a regulated environment, accountability sits with the organisation that deploys the agent, even where the underlying system was built by a third party.

The two questions from the data literacy and governance thread are more important at this level than anywhere else. Before handing a process to an agent, be certain you have answered both.

6a Human-in-the-loop

The system plans and executes but pauses for approval at defined checkpoints. For most finance use cases today, this is the safer default because it aligns better with control, audit, and exception-management requirements.

6b Fully autonomous

The system executes without requiring approval at each step, intervening only when a defined threshold is breached. This is powerful, and increasingly where major platforms are heading, but it is a controls decision before it is a technology decision.

One important point to watch is that the practical meaning of “adequate oversight” is evolving. As agentic systems scale, continuous monitoring and the ability to intervene quickly may become more realistic than pausing for approval at every step. That does not change where most finance use cases should sit today, but it does mean 6a and 6b are better understood as points on a spectrum than as a permanent binary choice.

Level 6 (continued)

A NOTE ON COPILOT, STUDIO, AND COWORK

Standard Copilot in everyday Microsoft 365 apps sits closer to Level 3 - it responds when asked in a single turn and does not act on its own. Copilot Studio is not a fixed level. It is a build environment, and where a particular solution sits depends entirely on how it is configured.

A purely conversational agent - for example, an internal expense-policy assistant that answers questions and waits to be asked - may sit at Level 5. If the same Studio-built agent is given an autonomous trigger, such as monitoring submitted expense claims overnight and flagging anomalies before anyone asks it to, it moves into Level 6 because it is now deciding what to do next on its own.

Copilot Cowork is designed for agentic, multi-step work, but it should not be assumed to sit at a single maturity level in every use case. Microsoft's own guidance for building agents points to approval checkpoints for sensitive actions, so a Copilot Studio agent or a Cowork deployment can each be configured at 6a or 6b depending on the specific process, permissions, and thresholds involved.

Do not assume "we use Copilot" tells you where you sit. Ask whether you are talking about the everyday assistant, something built in Studio, or Cowork; and if it is Studio-built, ask what triggers it and what it has been set up to do without human intervention.

TRY THIS

Do not start here by handing over something consequential. Pick a low-stakes task with a clear, checkable outcome, and watch how the system breaks the task down before you trust it with more.

If your organisation has access to Copilot Cowork or similar, Microsoft's own [Copilot Studio guidance on designing autonomous agents](#) is a useful read before deploying anything that runs unwatched, specifically on where to place human approval checkpoints.

If you are starting from scratch with Copilot Studio itself, Microsoft runs a free [Copilot Studio Agent Academy](#) with self-paced material, plus regular free one-day live courses - worth checking directly, since the specific programme names and links tend to change.

Where to go from here

Knowing where you sit on this ladder is the start, not the destination. The real value comes from applying it to your own finance function: which level is your team operating at, where are the gaps, and what would move the needle first.

Beyond Level 6

It is also worth knowing where this ladder deliberately stops. Beyond Level 6 sits software engineering specialism proper - building and maintaining agent infrastructure, custom model deployments, and the like. That is a genuine IT or engineering discipline, not a finance-function skill, and it belongs with the team who owns it. Recognising that boundary is as useful as knowing where you sit within it.

About Cathy Brown Advisory

That is also where this document ends and the actual work begins. I fix and automate finance functions for UK SMEs and mid-market businesses. Manual spreadsheets, slow month-ends, and missing forecasts are common symptoms, not an exhaustive list, of a finance function that has outgrown its processes - and I help rebuild it into one that runs itself.

Finance specialist

I am a chartered accountant with over 20 years in senior finance, including group reporting across more than 100 entities, and I pair that with hands-on AI and automation delivery. That combination is the point: I have run the functions this framework describes, not just automated around the edges of one I do not understand.

Independent

Unlike many agencies, I have no software partnerships or reseller commissions steering the recommendation. The roadmap reflects what your function needs, not what I am incentivised to sell.

Data governance and compliance

The ACA qualification means data security and compliance risk are not an afterthought - they are second nature. Client confidentiality, audit trails, and regulatory obligations are things I have been trained and examined on, and they shape how I approach every recommendation, not something bolted on once something goes wrong.

The initial offer

In practice, that starts with a free questionnaire, which gives you an initial assessment of your AI and automation readiness. Should you wish to explore further, the next step is a full AI and automation readiness assessment, involving a roughly hour-long call. The output is a prioritised roadmap - the specific improvements with the highest return, in the order worth tackling them.

Cathy Brown, ACA

Founder, Cathy Brown Advisory
cathybrownadvisory.co.uk